

---

# Linear filtering

1

---

# Roadmap

Machine Vision Technology							
Semantic information				Metric 3D information			
Pixels	Segments	Images	Videos	Camera		Multi-view Geometry	
<b>Convolutions</b> Edges & Fitting Local features Texture	Segmentation Clustering	Recognition Detection	Motion Tracking	Camera Model	Camera Calibration	Epipolar Geometry	SFM

2

## Types of Images

Binary



Gray Scale




Color



Source: Ulas Bagci

3

## Binary image representation




Source: Ulas Bagci

4

## Grayscale image representation

Slide credit:  
Ulas Bagci



5

## Color Image - one channel



Source: Ulas Bagci

6

## Color image representation

---



Source: Ulas Bagci

7

## Motivation: Image denoising

---

- How can we reduce noise in a photograph?



Source: S. Lazebnik

8

## Moving average

---

- Let's replace each pixel with a *weighted* average of its neighborhood
- The weights are called the *filter kernel*
- What are the weights for the average of a 3x3 neighborhood?

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

"box filter"

Source: S. Lazebnik

9

## Defining convolution


---

- Let  $f$  be the image and  $g$  be the kernel. The output of convolving  $f$  with  $g$  is denoted  $f * g$ .

$$(f * g)[m, n] = \sum_{k, l} f[m - k, n - l] g[k, l]$$

b	u	!
p	e	j
a	q	c

Convention:  
kernel is "flipped"



10

## Key properties

---

- **Linearity:**  $\text{filter}(f_1 + f_2) = \text{filter}(f_1) + \text{filter}(f_2)$
- **Shift invariance:** same behavior regardless of pixel location:  $\text{filter}(\text{shift}(f)) = \text{shift}(\text{filter}(f))$
- Theoretical result: any linear shift-invariant operator can be represented as a convolution

Source: S. Lazebnik

11

## Properties in more detail

---

- **Commutative:**  $a * b = b * a$ 
  - Conceptually no difference between filter and signal
- **Associative:**  $a * (b * c) = (a * b) * c$ 
  - Often apply several filters one after another:  $((a * b_1) * b_2) * b_3$
  - This is equivalent to applying one filter:  $a * (b_1 * b_2 * b_3)$
- **Distributes over addition:**  $a * (b + c) = (a * b) + (a * c)$
- **Scalars factor out:**  $ka * b = a * kb = k(a * b)$
- **Identity:** unit impulse  $e = [\dots, 0, 0, 1, 0, 0, \dots]$ ,  
 $a * e = a$


Source: S. Lazebnik

12

## Annoying details

What is the size of the output?

- MATLAB: `filter2(g, f, shape)`
  - `shape = 'full'`: output size is sum of sizes of `f` and `g`
  - `shape = 'same'`: output size is same as `f`
  - `shape = 'valid'`: output size is difference of sizes of `f` and `g`



Source: S. Lazebnik

13

## Annoying details

What about near the edge?

- the filter window falls off the edge of the image
- need to extrapolate
- methods:
  - clip filter (black)
  - wrap around
  - copy edge
  - reflect across edge



Source: S. Marschner

14

## Annoying details

---

### What about near the edge?

- the filter window falls off the edge of the image
- need to extrapolate
- methods (MATLAB):
  - clip filter (black): `imfilter(f, g, 0)`
  - wrap around: `imfilter(f, g, 'circular')`
  - copy edge: `imfilter(f, g, 'replicate')`
  - reflect across edge: `imfilter(f, g, 'symmetric')`

Source: S. Marschner

15

## Practice with linear filters

---



Original

0	0	0
0	1	0
0	0	0


?

Source: D. Lowe

16


## Practice with linear filters

---



Original

0	0	0
0	1	0
0	0	0



Filtered  
(no change)

Source: D. Lowe

17

## Practice with linear filters

---



Original

0	0	0
0	0	1
0	0	0

?

Source: D. Lowe

18


## Practice with linear filters

---



Original

0	0	0
0	0	1
0	0	0



Shifted *left*  
By 1 pixel

Source: D. Lowe

19

## Practice with linear filters

---



Original

 $\frac{1}{9}$ 

1	1	1
1	1	1
1	1	1

?

Source: D. Lowe

20


## Practice with linear filters

---



Original

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$




Blur (with a box filter)

Source: D. Lowe

21

## Practice with linear filters

---



Original

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

?

(Note that filter sums to 1)

Source: D. Lowe

22

## Practice with linear filters

---



Original

0	0	0
0	2	0
0	0	0

-  $\frac{1}{9}$

1	1	1
1	1	1
1	1	1



### Sharpening filter


- Accentuates differences  
with local average

Source: D. Lowe


23

## Sharpening

---



before



after

Source: D. Lowe

24

## Sharpening

What does blurring take away?




Source: D. Lowe

25

## Sharpening

What does blurring take away?



Let's add it back:




Source: D. Lowe

26

## Smoothing with box filter revisited

---

- What's wrong with this picture?
- What's the solution?




Source: D. Forsyth

27

## Smoothing with box filter revisited

---

- What's wrong with this picture?
- What's the solution?
  - To eliminate edge effects, weight contribution of neighborhood pixels according to their closeness to the center





"fuzzy blob"

Source: S. Lazebnik

28

## Gaussian Kernel

$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$



0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

5 x 5,  $\sigma = 1$


- Constant factor at front makes volume sum to 1 (can be ignored when computing the filter values, as we should renormalize weights to sum to 1 in any case)

Source: C. Rasmussen


29

## Gaussian Kernel

$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$



$\sigma = 2$  with 30 x 30 kernel



$\sigma = 5$  with 30 x 30 kernel


- Standard deviation  $\sigma$ : determines extent of smoothing

Source: K. Grauman

30

## Choosing kernel width

- The Gaussian function has infinite support, but discrete filters use finite kernels




Source: K. Grauman

31

## Choosing kernel width

- Rule of thumb: set filter half-width to about  $3\sigma$




Source: S. Lazebnik

32

## Gaussian vs. box filtering

---



Source: S. Lazebnik

33

## Gaussian filters

---

- Remove “high-frequency” components from the image (low-pass filter)
- Convolution with self is another Gaussian
  - So can smooth with small- $\sigma$  kernel, repeat, and get same result as larger- $\sigma$  kernel would have
  - Convoluting two times with Gaussian kernel with std. dev.  $\sigma$  is same as convoluting once with kernel with std. dev.  $\sigma\sqrt{2}$
- *Separable* kernel
  - Factors into product of two 1D Gaussians

Source: K. Grauman

34

## Separability of the Gaussian filter

$$\begin{aligned}
 G_{\sigma}(x, y) &= \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right) \\
 &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{x^2}{2\sigma^2}\right)\right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{y^2}{2\sigma^2}\right)\right)
 \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of  $x$  and the other a function of  $y$

In this case, the two functions are the (identical) 1D Gaussian

Source: D. Lowe

35

## Separability example

2D convolution (center location only)

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array}$$

The filter factors into a product of 1D filters:

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array}$$

Perform convolution along rows:

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array} =$$

Followed by convolution along the remaining column:

Source: K. Grauman

36

## Why is separability useful?

- What is the complexity of filtering an  $n \times n$  image with an  $m \times m$  kernel?
  - $O(n^2 m^2)$
- What if the kernel is separable?
  - $O(n^2 m)$

Source: S. Lazebnik

37

## Noise



Original



Salt and pepper noise



Impulse noise



Gaussian noise


- **Salt and pepper noise:** contains random occurrences of black and white pixels
- **Impulse noise:** contains random occurrences of white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian normal distribution

Source: S. Seitz

38

## Gaussian noise


- Mathematical model: sum of many independent factors
- Good for small standard deviations
- Assumption: independent, zero-mean noise



Source: M. Hebert

39

## Reducing Gaussian noise



Smoothing with larger standard deviations suppresses noise, but also blurs the image

Source: S. Lazebnik

40

## Reducing salt-and-pepper noise




What's wrong with the results?

Source: S. Lazebnik

41

## Alternative idea: Median filtering

- A **median filter** operates over a window by selecting the median intensity in the window



- Is median filtering linear?


Source: K. Grauman

42

## Median filter

- What advantage does median filtering have over Gaussian filtering?
  - Robustness to outliers

filters have width 5 :




Source: K. Grauman



43

## Median filter

Salt-and-pepper noise




Median filtered



Source: M. Hebert

44


## Gaussian vs. median filtering



Source: S. Lazebnik

45

## Sharpening revisited



Source: D. Lowe

46

## Sharpening revisited

What does blurring take away?




Let's add it back:



Source: S. Lazebnik

47

## Unsharp mask filter



Source: S. Lazebnik

48